



Data Dissemination

Contents:



- Some Apps
- Nature of information
- Modes of interaction
- Data delivery options
- MOM
 - Addressing options

Mariano Cilia / mcilia@gmail.com

1



Motivated by Examples

- Traffic management
 - Air-traffic control
 - Emergency vehicles
- Environmental
 - Concentration of chemicals
 - Dangerous emissions
 - Forest fires
 - Flood control

2



Example – Flood Control

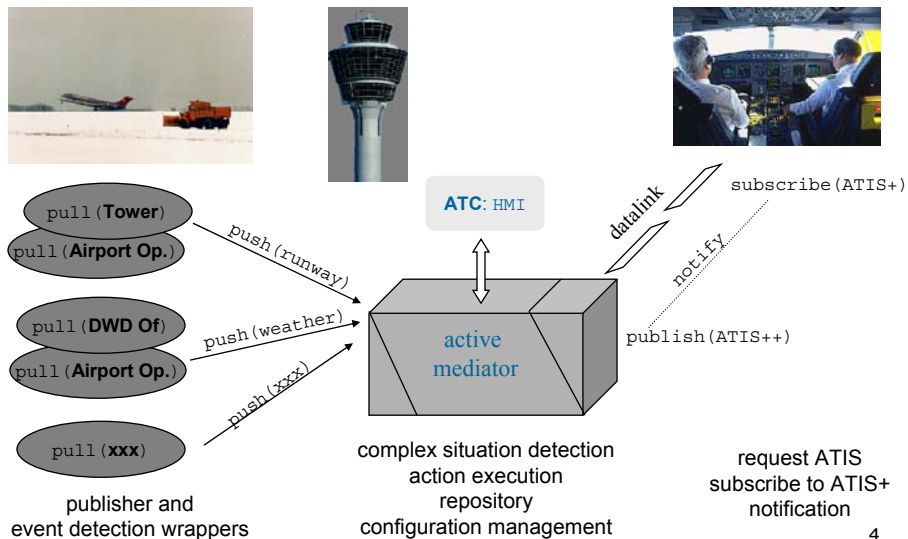
- A combination of happenings (coming from sensor data) determines a situation
- Polling sensors
 - Too frequent
 - Consumption of resources
 - Too infrequent
 - Malfunction



3



Application scenario: ATIS



4



Examples (cont.)

- Plant and reactor control
 - Equipment control
- Defense
 - Missile detection
 - Battlefield monitoring
- Workflow management

5



Examples (cont.)

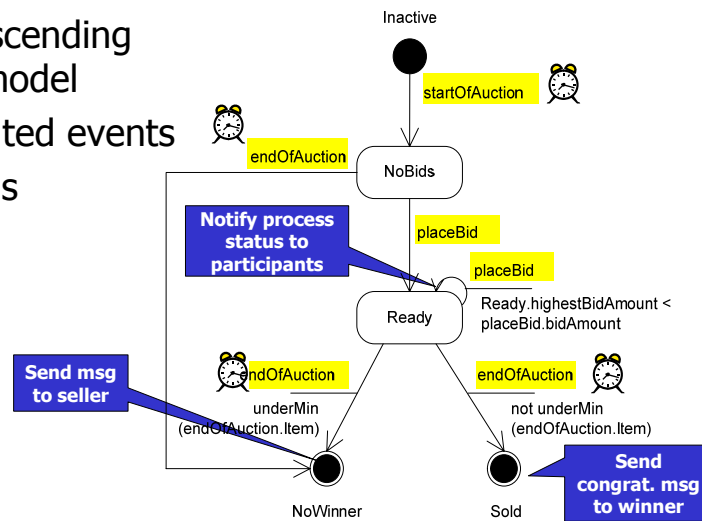
- Commerce
 - Inventory control
 - Supply Chain Management
 - Marketplaces
 - e-Auctions
 - Online shops

6

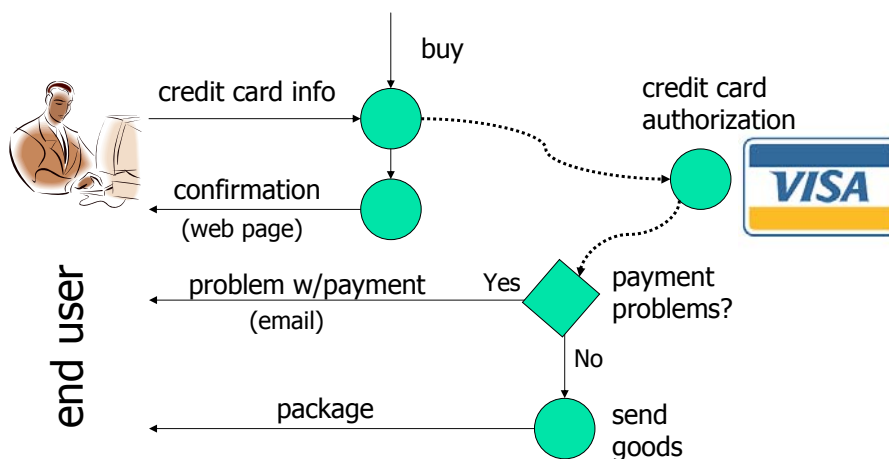


Example – e-Auctions

- Simple ascending auction model
- Time-related events
- Conditions



Example – Online shops





RFID – Supply Chain Mgmt (cont)

HOW EFFICIENT IS YOUR SUPPLY CHAIN?

XPLANATIONS® by XPLANE®


The EPC™ Network can track individual items throughout the supply chain, from manufacture to sale. This will revolutionize the way people buy, sell and distribute products. Here's how it works.

1. ON THE RETAIL FLOOR

The moment a customer takes a product from the shelf, "smart shelves" automatically order more. Stock people and distributors keep the shelves full so customers can buy what they want.


AND IN THE BACK ROOM

The EPC™ Network tells the retailer exactly what's on the shelf and in the stock room as well as what's rolling off the truck. There is no need for clerks to maintain costly buffer stock or manually break a pallet down in order to check every case.



HOW IT WORKS

Every item contains a microchip with a unique identifier — called an Electronic Product Code (EPC™). This Radio Frequency Identification (RFID) tag allows precise tracking of the product. Cases and pallets can carry their own unique tags. Learn more about the EPC™ Network at the Auto-ID Center's web site at autoidcenter.org.




2. IN THE WAREHOUSE

The warehouse manager quickly routes shipments to the right place because he can look up what's in the warehouse and on every truck. Duplicate and missed shipments are a thing of the past.

3. IN THE PLANT

Manufacturers organize their runs based on up-to-date information and know an item's destination as it comes off the line. If an incident involving a defect or tampering arises, only the affected products need to be recalled.

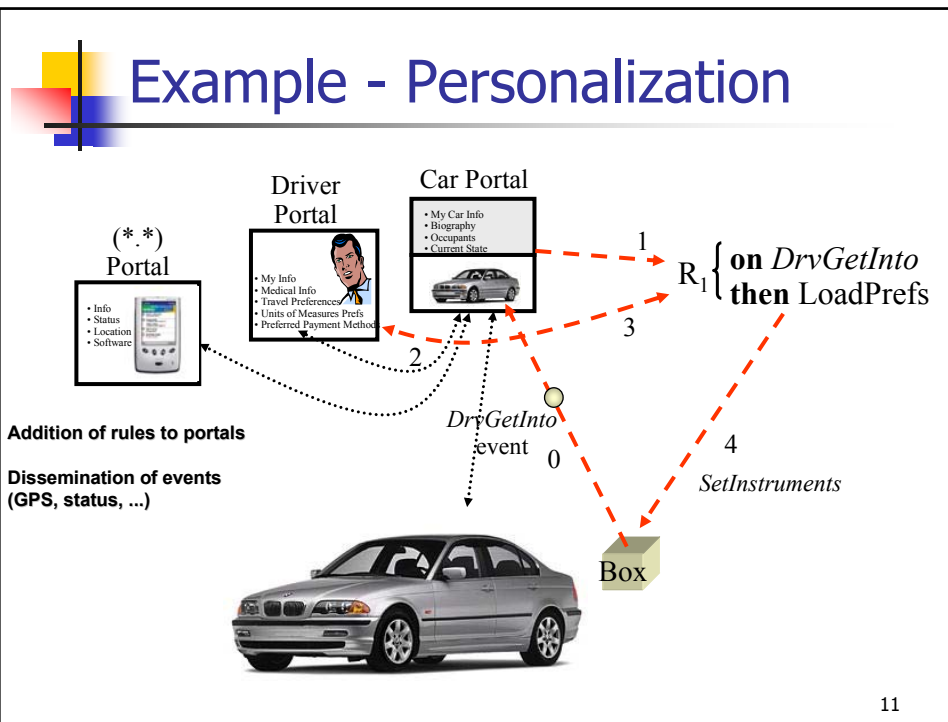


The Auto-ID Center | © 2000 XPLANE, Inc.®



Examples (cont.)

- Personalization
 - User Interfaces
 - Services



- ## Examples (cont.)
- Personalization
 - User Interfaces
 - Services
 - Financial applications
 - Commodity trading
 - Currency trading
 - Stock trading
- 12



Example – Stock trading

- Sample
 - **ON** stock.Name=IBM
 - **IF** stock.Price<20
 - **THEN** call myBuy()
- High volume



13



Convergence of Technologies

- Ambient Intelligence and smart devices require continuous monitoring of events
- Miniaturization of sensors, ubiquitous deployment
- Context information for proper interpretation
- (Almost) complete reachability of individuals causes unbounded appetite for information
- Need to filter and interpret large amounts of heterogeneous and short-lived data
- Large distributed systems must detect and correct failures/exceptions (autonomic computing, ESCM, zero latency enterprise)

14



The Nature of Information



- Information *flows* from producer to consumer
 - info-pipes, broadcast disks, event streams, pub/sub
- Static view of information is a simplification
 - data flows into/out of high latency pool (database)

Mechanisms for access to static information (**queries**) are different from those for accessing flow of information (**subscription/filters**)

15



Working Hypothesis

- Desintermediation/Reintermediation
- B2C, B2B, B2B2C, C2C, Portals, Markets (DotComs vs. NewCos), m-commerce ...
- First generation e-commerce systems mapped existing applications 1:1 to new medium
- Next generation(s) will be based on flexible integration of services and components
- Flow of tasks and information
- How should (middleware) platforms look?

16



Modes of Interaction


		Initiator of Interaction	
		Consumer	Producer
Knowledge about Counterpart	Full	Request/Reply	One-to-One Message
	No	Anonymous Request/Reply	Event-based dissemination

1st generation

17




Modes of Interaction – R/R



18


Modes of Interaction – R/R



- Known server

19

Modes of Interaction – R/R



20



Request/Reply

- Direct and synchronous communications
 - Enforces tightly coupling of comm. parties
 - Impairs scalability
- Clients pull remote data sources
 - Trade off
 - Usage of data vs. data accuracy
 - Short polling interval → waste resource
 - Long polling interval → increase update latency
- Need for asynchronous and decoupled operations

21



Request/Reply (cont.)

- Simple
 - + imperative nature of C/S paradigm
 - + programming language abstraction
- Drawbacks
 - Point-to-point communication limits scalability
 - Polling limits accuracy of data
 - Unnecessary bandwidth consumption

22



Modes of Interaction

		Initiator of Interaction	
		Consumer	Producer
Knowledge about Counterpart	Full	Request/Reply	One-to-One Message
	No	Anonymous Request/Reply	Event-based dissemination

23



Anonymous R/R

- Consumer does not specify the provider
- Request is delivered to an arbitrary set of providers
- Identity of provider is unknown
- Load balancing

24



Modes of Interaction

		Initiator of Interaction	
		Consumer	Producer
Knowledge about Counterpart	Full	Request/Reply	One-to-One Message
	No	Anonymous Request/Reply	Event-based dissemination

25



1-to-1 message / Callback

- Consumer registers interest with a known provider
- Provider repeatedly evaluates interests
 - when true → callback registered consumer
 - Responsible for managing list of interests and registered consumers
- One to one message
- Observer pattern

26



Modes of Interaction


		Initiator of Interaction	
		Consumer	Producer
Knowledge about Counterpart	Full	Request/Reply	One-to-One Message
	No	Anonymous Request/Reply	Event-based dissemination

Needs a mediator

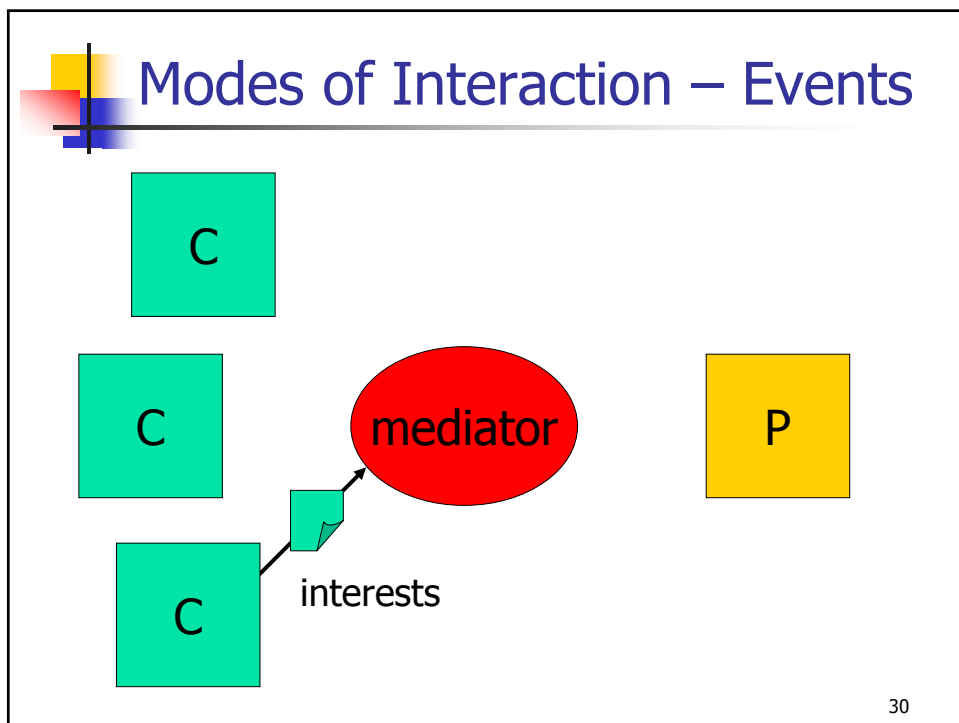
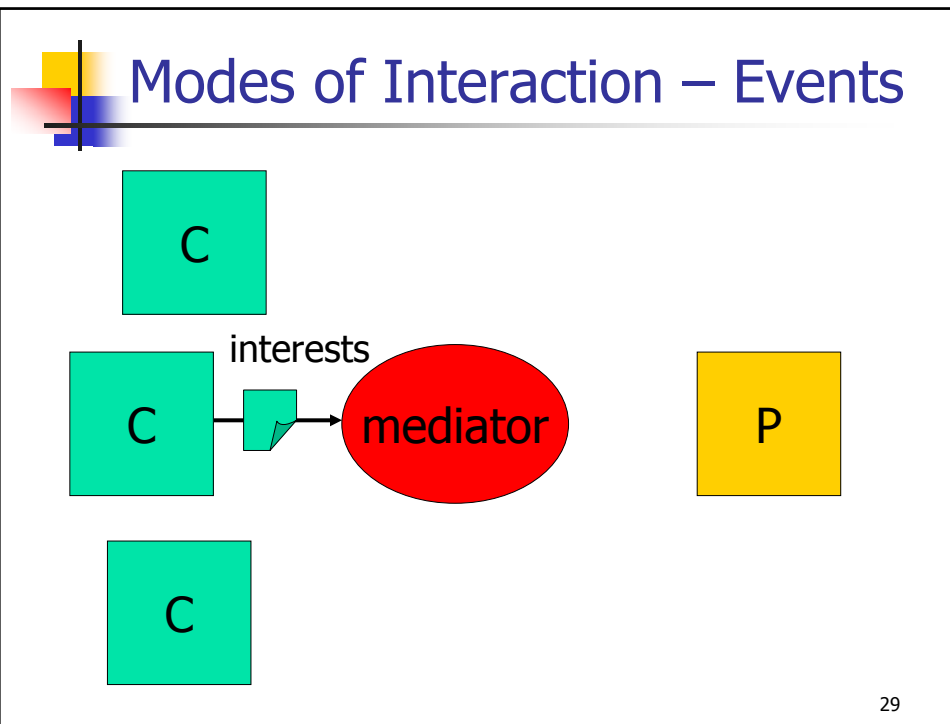
27

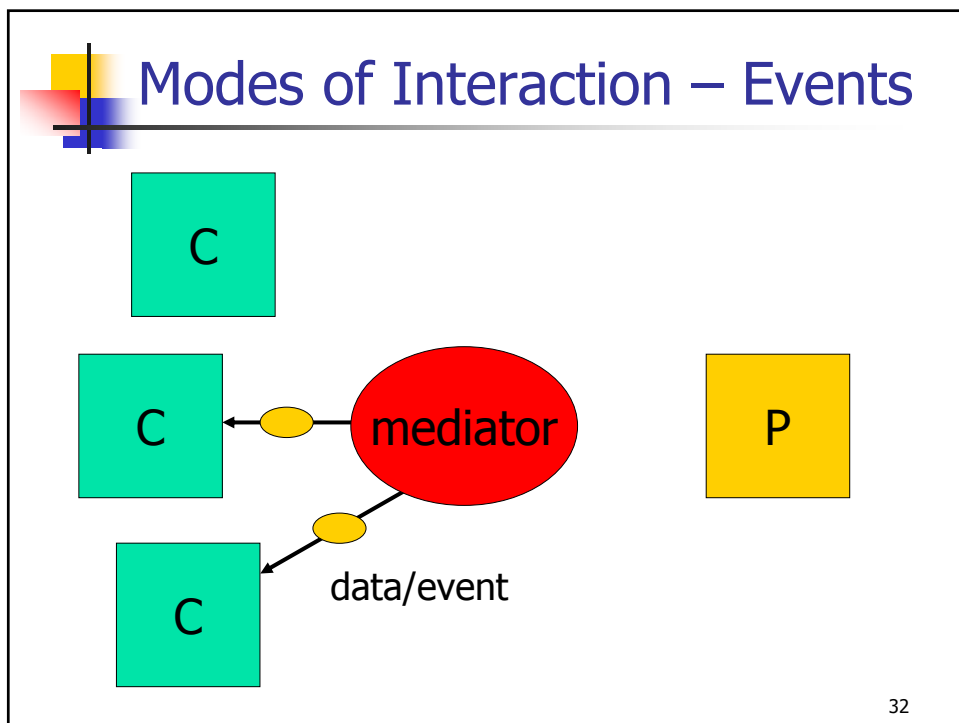
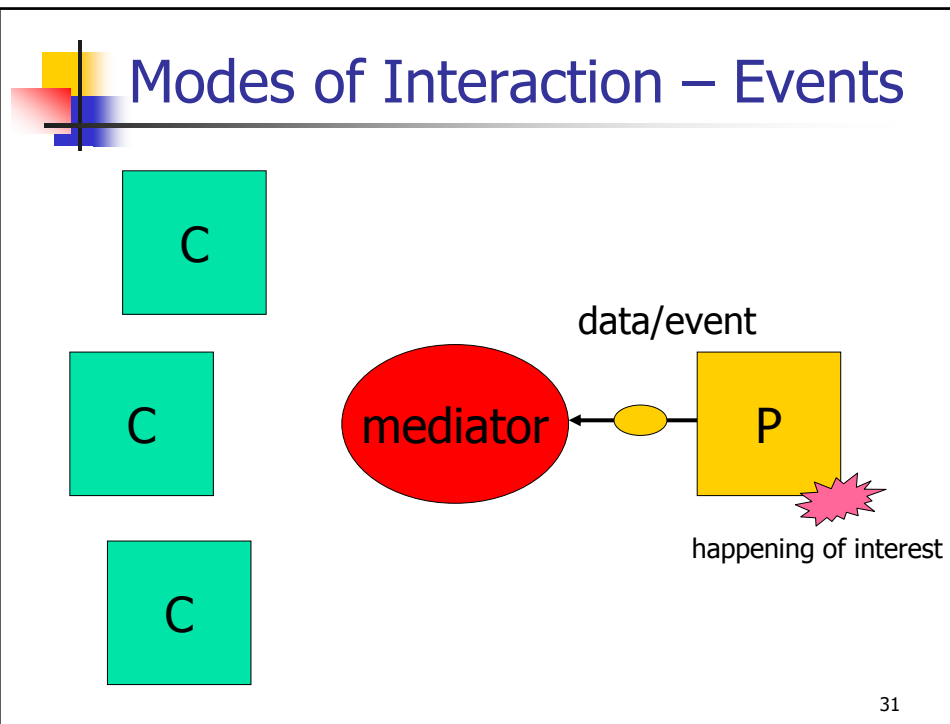


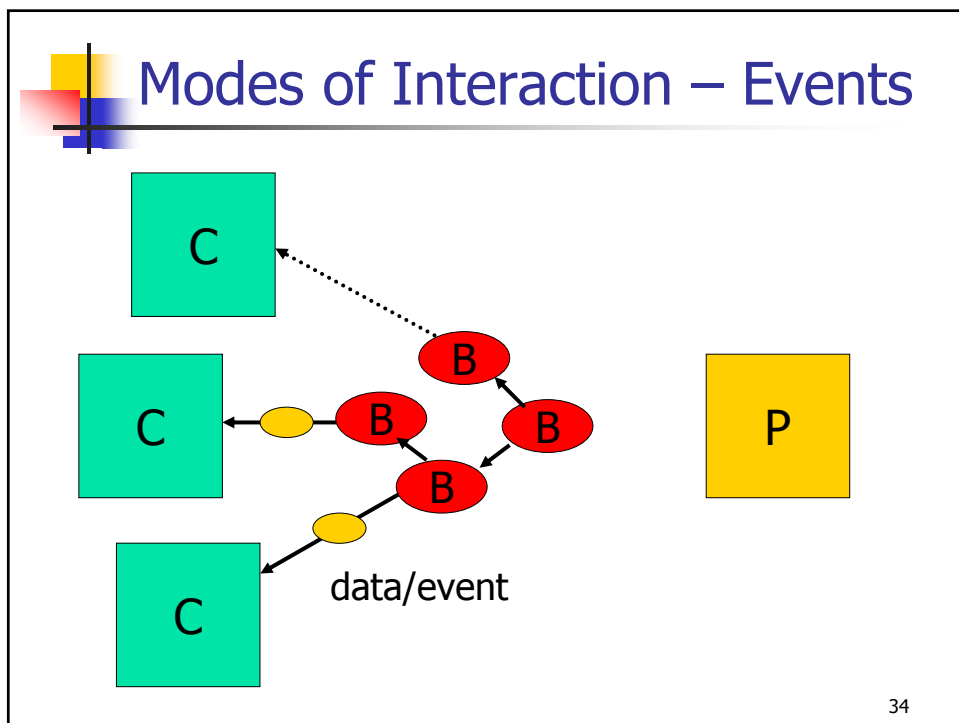
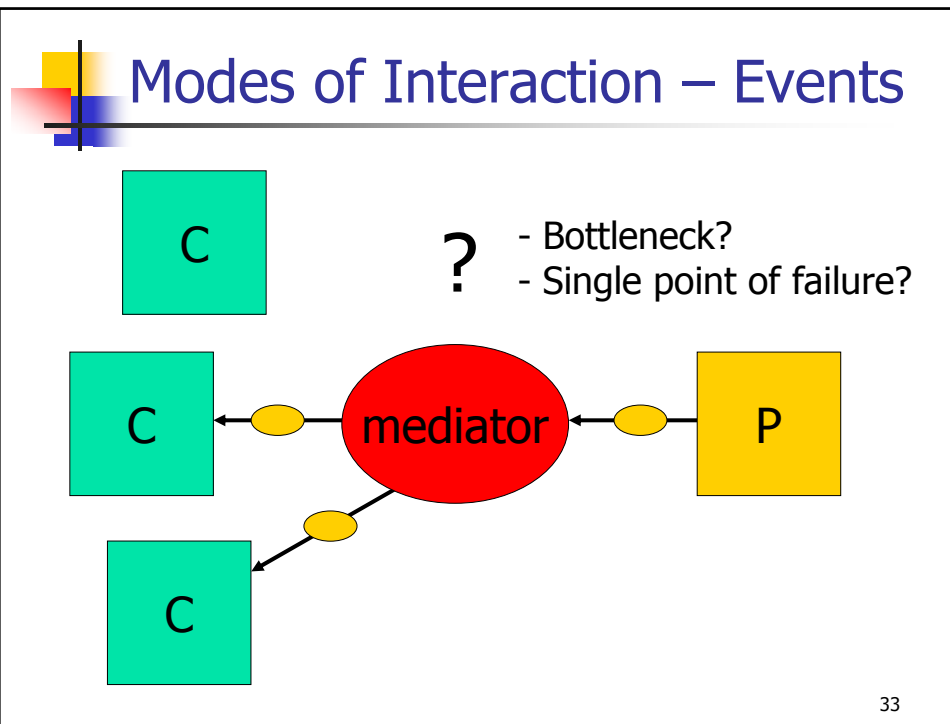
Modes of Interaction – Events



28









Events and Notifications

- Event
 - happening of interest at observed object(s)
- Notification
 - i) communication of event occurrence to interested recipients
 - ii) reification of observed event
- Notification Service (NS)
 - provides infrastructure to register for and deliver notifications
 - i.e. publish(), subscribe(), notify()

35



Event Notification - Patterns

- Observer: observable events
 - Event producers have knowledge about event consumers
- Mediator: centralized mediation
 - Encapsulates and coordinates communication
- Notification Service
 - Combines the Observer and the Mediator patterns
 - Subscribers only know about events not about publishers
 - Mediation between event producers and consumers

36



Event-based Paradigm

- The (data) **producer is the initiator** of communication
- Notifications are **not addressed** to any specific consumer
 - Producers are not aware of consumers
- Consumers issue **subscriptions** (interests)
 - Consumers are not aware of producers
- Notifications are delivered to consumers if they match with subscriptions
- Flexible!

37



Modes of Interaction

- Influences
 - The architecture of the system
 - The design of the individual processes involved
- Link distributed parts of the system
 - difficult to change afterwards
- MoI determine system's ability to adapt, evolve and scale
- MoI is confused with the implementation techniques

38



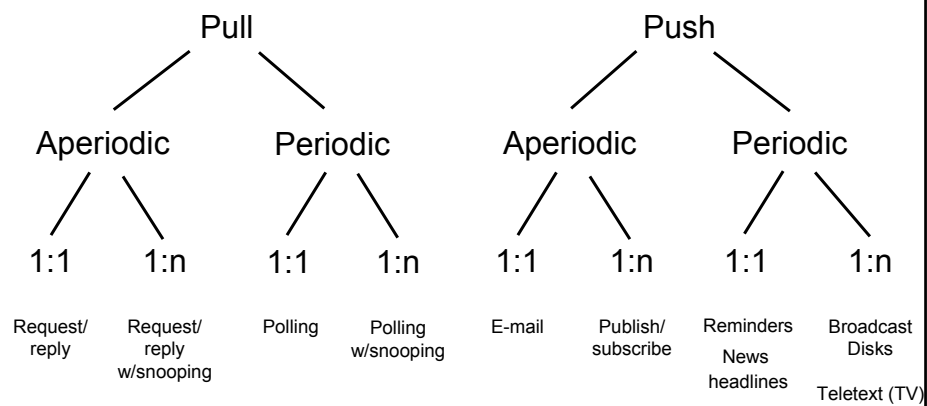
Interaction Patterns

- Initiation
 - (client) pull vs. (server) push
 - periodic vs. aperiodic
- Topology
 - 1:1 (unicast) vs. 1:n (multicast)
- Lifecycle
 - time-dependent vs. time-independent
- Concurrency
 - blocking vs. non-blocking
- Reliability
 - atomic, at-least-once, at-most-once, exactly-once

39



Data Delivery Options



40



Interaction vs. Invocation

- Must separate mode of interaction and implementation (invocation) technique
- Separation must occur at various levels of abstraction
 - RPC implemented using messages
 - Implementation using other interaction patterns
 - Pointcast: implemented an event-driven notification service through a polling mechanism

41



Invocation Mechanisms of C/S Sys

- The communication mechanisms used in client/server systems fall into one of the following categories:
 - remote procedure call (RPC)
 - transactional RPC
 - peer-to-peer messaging
 - queues
 - transactional queues
 - events/Publish-Subscribe

42



Middleware

- used to glue together applications (components):
 - IPC by sockets, shared memory
 - TCP/IP, X.25
 - common database
 - RPC, CORBA RMI, J2EE
 - MOM

43



Message Oriented Middleware

- applications communicate through explicitly sending/receiving messages
- most common flavors:
 - queues
 - point-to-point (mostly)
 - location-based addressing
 - enqueue, dequeue
 - store and forward
 - publish/subscribe
 - different addressing approaches
 - register & callback (Observer pattern)
 - optimize network use

44



MOM (cont.)

- flexible interaction
 - C/S request/reply, one-way push
 - asynchronous and time-independent
 - 1:1, n:1, 1:n, m:n
 - priorities
- flexible reliability
 - volatile/persistent/transactional queues
 - reliable/certified/transactional pub/sub
- additional services
 - load balancing, naming, security, content transformation

45



Communication Mechanisms



Already seen Request/Reply

- Simple
 - + imperative nature of C/S paradigm
 - + programming language abstraction
- Drawbacks
 - Point-to-point communication limits scalability
 - Polling limits accuracy of data
 - Unnecessary bandwidth consumption

47



Queues

- Why use queues?
 - Asynchronous communication
 - No blocking while waiting for reply
 - Clients can submit requests even if server is not available
 - Easy to handle results of disconnected clients
 - Load balancing
 - Possibility of prioritizing the requests in the queue

48



Operation with Queues

- Persistent queue between client and server
- Client: enqueues requests, dequeues replies
- Server: dequeues a request, processes request, enqueues reply, commits
- If transaction aborts due to system reasons it is enqueued again

49



Queue Managers

- Queue manager needed
 - operations on queue elements: enqueue, dequeue, scan queue, keyed access
 - create and destroy queues
 - modify a queue's attributes, such as owner, size, privileges
 - start and stop queue
 - routing of requests (forwarding to another queue manager in case of overload)

50



Server's View of Queuing

- Assume each request is for execution of a single transaction
- Server dequeues a request, executes the request, enques the result, and commits
- If the transaction aborts
 - the dequeue operation is undone
 - the enqueue operation is undone if already started
- If client checks queues, request is either in request queue, in process, or result in result queue

51



Client's View of Queuing

- Client perceives three transactions for each request:
 - one transaction to enqueue request
 - receive input from user, construct request, enqueue request, commit
 - one server transaction (described before)
 - one transaction to dequeue results
 - dequeue reply from result queue, convert to proper output format, deliver output, commit (wiping out result in result queue)

52


Request/Reply with Queues

Client A

TX 1:
Start
get input
construct request
enqueue request
Commit

TX 3:
Start
dequeue reply
decode reply
process output
Commit

Message queues are good for
asynchronous point to point
(1:1) messaging



Server B

TX 2:
Start
dequeue request
process request
enqueue reply
Commit

54

Cost/Benefit of Operating with Queues

- Using queues buys flexibility
 - communication with unavailable clients or servers
 - load balancing across servers
 - easy implementation of priorities
 - easier integration of legacy systems
- Using queues is expensive
 - 3 transactions instead of one
 - transactional queues must be managed by a (specialized) DBMS to guarantee persistence and transaction semantics

55



Need for Persistent Sessions

- System must be able to identify sending and receiving transactions and match them
- Without request/reply semantics, queue manager may not accept requests with output parameters (since results would be simply dumped on device)
- Recovery of queuing systems later (with TPM)



56



Summary communication mechanisms

- RPC: synchronous, simple call-return semantics, hard-wired termination and ordering
- Multicast: 1:N messaging for group communication
- Queues: fully asynchronous, maximum flexibility for handling client/server/communication failures

58




Pub/Sub Notification Service

- Main characteristics
 - decouples producers and consumers
 - anonymous to each other
 - dynamic number of consumers and producers
 - no directory service is needed
- Addressing models
 - Channel-based
 - Subject-based
 - Content-based
 - Concept-based

60

Channel-based



- Less powerful
- Simple

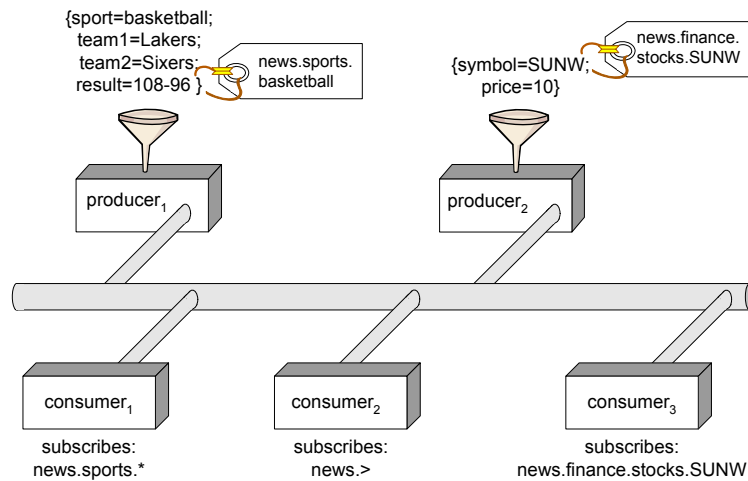
61

Subject-based Addressing

- Subject-based addressing avoids use of physical network addresses
- Senders label a data message with a subject name
 - Subject = characterize/synthesize message content
- Consumers listen to names and pick up the messages with the proper subject name
- Anonymous rendezvous:
 - producers need not know how data is consumed
 - consumers need not know how or where data is produced

62

Subject-based Addressing (2)



63

Subject-based Addressing (3)

- Agreement on subject names
- New subjects can be created dynamically
- Subject names consist of elements (subject name hierarchy)
 - element.subelement.subsubelement
- Wildcards can be used
 - `RUN.*` matches `RUN.AWAY`
 - `RUN.*` matches `RUN.home`
 - `RUN.>` matches `RUN.AWAY.far`
- Difficult to change subject hierarchy

64



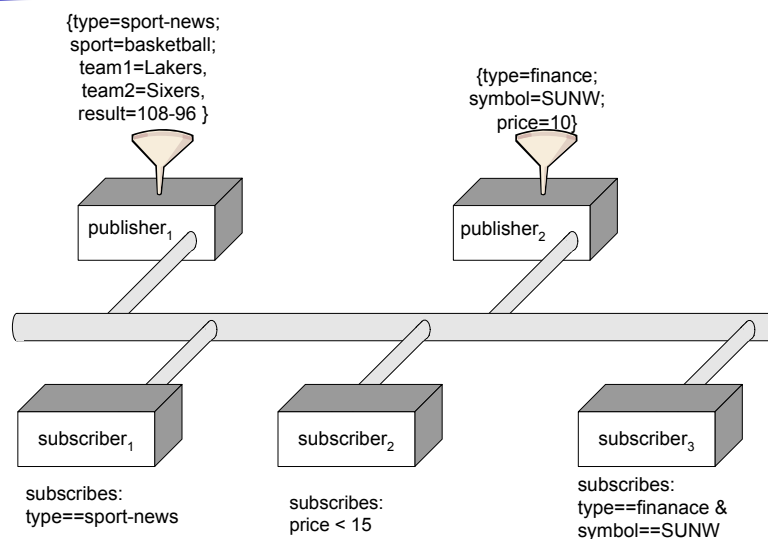
Content-Based Pub/Sub

- A content-based filter F
 - is a predicate on the content of notifications
 - induces the set of matching notifications
- Content-Based filtering is flexible but complex
 - cannot be easily mapped to "IP-Multicast"
- Centralized implementations not scalable to wide-area scenario
 - powerful distributed infrastructure required

65



Content-Based Pub/Sub



66



Problems derived from scale


- Flooding of notifications is not an applicable solution
 - need strategies for filter placement to optimize bandwidth and size of routing tables

67



Content-Based Routing

- Cooperating brokers
 - Local clients
 - Notification forwarding
 - Filter-Based Routing Tables
- Tradeoff: Network resource waste vs. filtering overhead (processing and delay)




subscriptions: F and G

68



Content-Based Routing

- Cooperating brokers
 - Local clients
 - Notification forwarding
 - Filter-Based Routing Tables
- Tradeoff: Network resource waste vs. filtering overhead (processing and delay)




routing tables are built

69



Content-Based Routing

- Cooperating brokers
 - Local clients
 - Notification forwarding
 - Filter-Based Routing Tables
- Tradeoff: Network resource waste vs. filtering overhead (processing and delay)




notification n is published
it matches with F

70

Content-Based Routing

- Cooperating brokers
 - Local clients
 - Notification forwarding
 - Filter-Based Routing Tables
- Tradeoff: Network resource waste vs. filtering overhead (processing and delay)



notification m is published
it matches with G

71

Content-Based Routing (cont)


- Size of routing tables crucial for scalability
 - global knowledge about all active subscriptions not feasible
- Solution: reduce size of routing tables and overhead to update them by
 - exploiting similarities among filters
 - identity tests
 - covering tests
 - merging of filters
 - trading accuracy vs. efficiency

72



Covering

- Filters can cover each other
- Covering can decrease
 - size of routing tables
 - filter forwarding overhead





73



Merging of Filters

- Filters can be merged
 - *perfect*
 - *Imperfect*
- Merging generates new covers
 - similar benefits as covering



74



The REBECA Approach

- Prototype of notification infrastructure
(**REBECA Event-Based Electronic Commerce Architecture**)
(<http://www.gkec.informatik.tu-darmstadt.de/rebeca>)
- Content-based routing with optimizations
 - Flexible filter framework
 - Support for complex data types
- Structuring publish/subscribe systems
 - Scoping
 - Sessions

75



Concept-based Pub/Sub

- Main advantages of Publish/subscribe
 - decouples producers and consumers
 - anonymous to each other
- BUT even though consumer and producer use a common vocabulary (let's suppose this) assumptions of participants are implicit
 - (date) 7/11/2003 Which is the month?
 - (price) 200 Currency? €, U\$?...
- Subscriptions expressed on flat messages

76



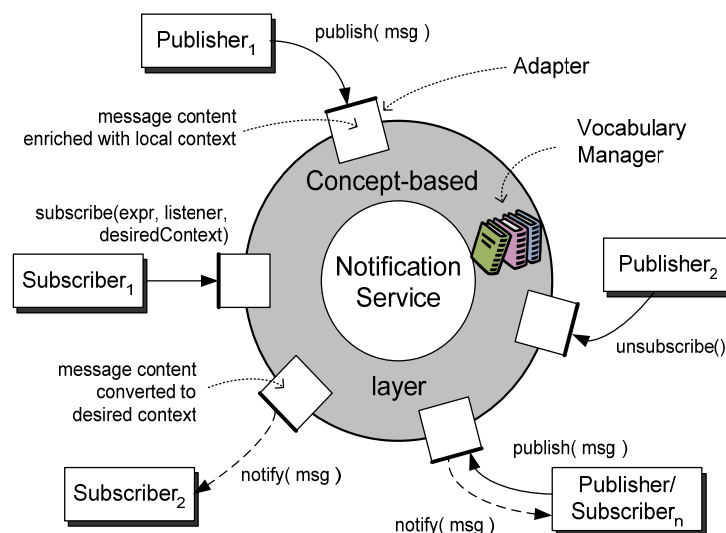
Concept-based Pub/Sub

- Provide a higher level of abstraction to describe the interests of publishers and subscribers
- Events represented using MIX
- Subscribers can specify their assumptions
 - Price < 100 [€]
 - DeliveryDate <= 7/11/2003 [dd/mm/yyyy]
- The notification service delivers **ready-to-process** events to subscribers
 - No further processing is needed

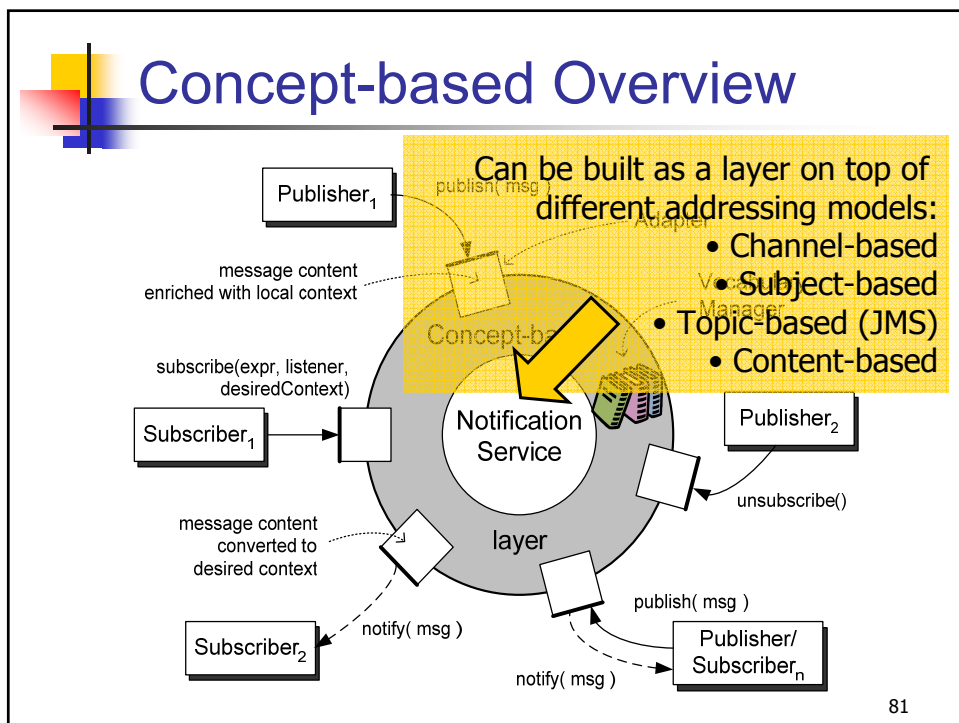
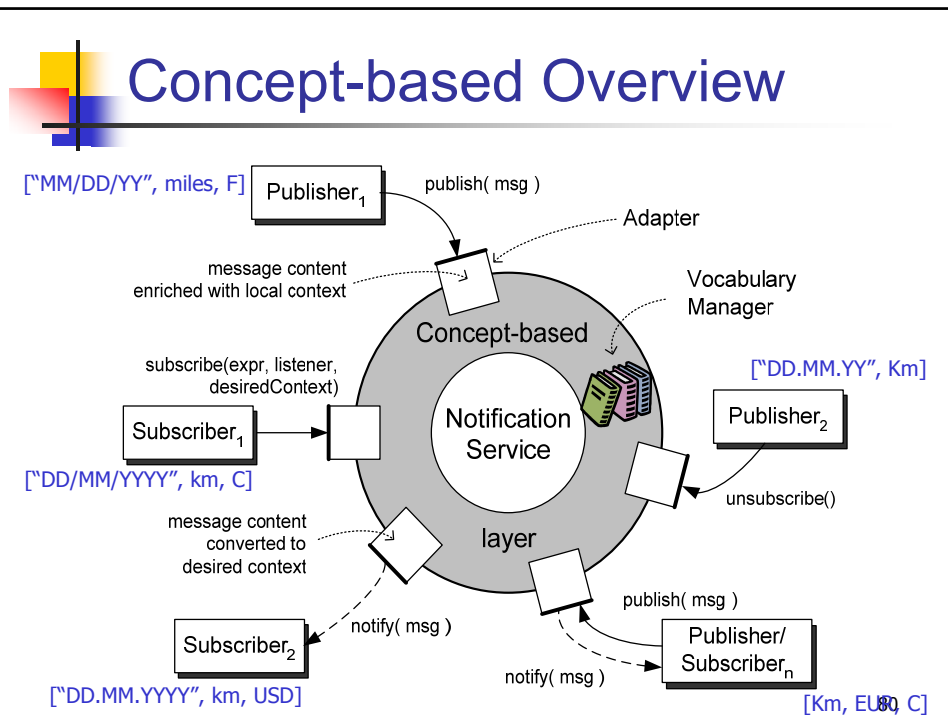
77



Concept-based Overview



79





Wrap up

- Different routing strategies according to application needs
- Filters/subscriptions on a single message
- Event correlation no supported
 - Need to cache/store semi-composed events
- Software Engineering
 - Need to scope events and subdivide event space

82



Data Dissemination Products



Queue Managers: IBM's MQSeries

- most TP Monitors offer queue managers (TUXEDO, Encina, TOP END)
- standalone products (IBM's MQSeries, BEA messageQ, SonicMQ, SUN JMQ, ...)
- MQSeries provides interoperable queue management across many Operating Systems
- works with all IBM TP monitors and any system supporting the X/Open XA interface (including CORBA OTS), Java connectivity included
- when working with a TPM, MQSeries uses the TPM transactions, otherwise it provides its own

92



MQSeries (cont.)

- multiple named queues supported
- queue forwarding among queues (e.g. for load balancing)
- queue forwarding occurs within channel agent's own transaction
- pub/sub brokering possible
- queue manager consists of
 - connection manager
 - data manager
 - lock manager
 - buffer manager
 - recovery manager
 - log manager

93



MQSeries: Qs, API ...

- types of Qs
 - local (app., transmission, dead-letter, initiation, ...)
 - remote, alias, model, dynamic
- interaction through MQI verbs
 - MQBEGIN, MQCMIT
 - MQPUT, MQGET (browsing, consuming, blocking/non-blocking)
 - control operations
 - connect/disconnect Qmanager (MQCONN, MQDISC)
 - set configurations, manage Q processing (MQOPEN, MQSET, MQCLOSE)
- interaction through C++/ Java APIs
- interaction through JMS API

94



MQSeries Messages

- messages can be
 - persistent
 - more secure, more expensive, logged, exactly once semantic
 - non-persistent
 - less secure, faster since in main memory, at most once semantic
- both types of messages can be enqueued in same queue
- message data
 - user defined format
 - default format and encodings

95



MQSeries: Messages (cont.)

- message consists of descriptor and data
- descriptor includes context
 - identity
 - origin
 - system message ID
 - application message ID
 - message type
 - datagram, request, reply,
 - report
- persistence flag
- name of destination queue
- ID of reply queue
- correlation ID
- expiry
- application-defined format
- report options
 - confirm on arrival, on delivery, on positive/negative action, on expiration, or on exception
- priority

96



MQSeries (cont.)

- management of message processing applications
 - process definition associated with Q
 - Qmanager sends trigger to *initiation* Q
 - trigger monitor may start application using process definition in trigger message

97



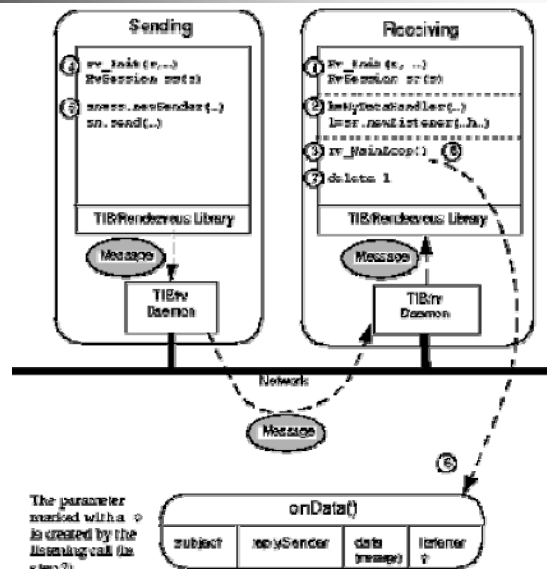
TIBCO's TIB/Rendezvous

- Event-driven, publish/subscribe
- Subject-based addressing
- Self-describing messages
- Leverage on broadcast & IP-multicast

98



TIB/Rendezvous Architecture



99



TIB/Rendezvous Messages

- Data Messages are self-describing
 - data + descriptive information
 - data
 - length of data
 - datatype indicator
 - subject name
 - listener callback functions receive same bundle
 - automatic conversion between local data format and TIB/Rendezvous wire format

100



Java Message Service - JMS

- Transactional, asynchronous messaging
- De-facto standard for Java messaging APIs
- Supported by almost every QueueManager vendor (IBM, Oracle, BEA,...)
- Many 100% Java, lightweight JMS products (Fiorano, Progress, Softwired, SpiritSoft, etc.)
- Designed for portability
 - Interfaces only => many different realizations
 - APIs to create, send, receive, read messages


105




JMS Model

- Supports both models

- Queues for point-to-point



- Publish/subscribe



106



JMS Model

- JMS point-to-point

- Queue object: encapsulates provider specific Q name
 - QueueConnection: handle to underlying transport
 - QueueSession: produces and consumes messages
 - TemporaryQueue: temporary storage for the QueueConnection
 - QueueConnectionFactory: creates QueueConnection
 - QueueReceiver: gets messages
 - QueueSender: puts messages

107



JMS pub/sub

- JMS publish/subscribe combines
 - Channels (known as topics) and
 - Expressions on envelope's attributes
- Factories, destinations, etc. identified via JNDI

108



JMS Messages

- Message types
 - text
 - map: (name,value) pairs
 - object: serializable object
 - stream: primitives
 - byte
- Message header used for addressing
- Message properties
 - list of (name,value) pairs
- Selectors are restricted to properties
 - SQL-like conditional expressions, `MyType='car'`
`AND MyName like 'Mu%'`


109



JMS Open Issues

- load balancing
- scalability/availability
- fault tolerance
- error notification
- end-to-end security
- segregation of domains
- simple and flexible deployment configuration
- Many of these issues being addressed by vendors

111



113